

kc_house_sales

October 4, 2022

0.1 Final Project Submission

Please fill out: * Student name: Emmanuel Kiplimo * Student pace: Full Time * Scheduled project review date/time: 3/10/2022 * Instructor name: William Okomba

0.2 Overview

A real estate agency in King County, Washington State is looking to explore sales and details on houses in the region to identify various features and factors that significantly contribute to the price of houses. As a Data Scientist, I'll analyze and model the data provided to draw insights and predictions on house prices.

The following queries will shed some light on house sales and features of an ideal house and a serve as build up for a predictive model:

1. Notable features that contribute to house prices.
2. Is there an observable trend of house sales across the year

0.2.1 Data understanding

```
[1]: # Import standard packages
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
/home/kiplimo/anaconda3/lib/python3.8/site-
packages/statsmodels/tsa/base/tsa_model.py:7: FutureWarning: pandas.Int64Index
is deprecated and will be removed from pandas in a future version. Use
pandas.Index with the appropriate dtype instead.
```

```
from pandas import (to_datetime, Int64Index, DatetimeIndex, Period,
/home/kiplimo/anaconda3/lib/python3.8/site-
packages/statsmodels/tsa/base/tsa_model.py:7: FutureWarning: pandas.Float64Index
```

is deprecated and will be removed from pandas in a future version. Use pandas.Index with the appropriate dtype instead.

```
from pandas import (to_datetime, Int64Index, DatetimeIndex, Period,
```

```
[2]: # Loading data onto a dataframe
df = pd.read_csv('data/kc_house_data.csv')
df.head()
```

```
[2]:
```

	id	date	price	bedrooms	bathrooms	sqft_living \
0	7129300520	10/13/2014	221900.0	3	1.00	1180
1	6414100192	12/9/2014	538000.0	3	2.25	2570
2	5631500400	2/25/2015	180000.0	2	1.00	770
3	2487200875	12/9/2014	604000.0	4	3.00	1960
4	1954400510	2/18/2015	510000.0	3	2.00	1680

	sqft_lot	floors	waterfront	view	...	grade	sqft_above \
0	5650	1.0	NaN	NONE	...	7 Average	1180
1	7242	2.0	NO	NONE	...	7 Average	2170
2	10000	1.0	NO	NONE	...	6 Low Average	770
3	5000	1.0	NO	NONE	...	7 Average	1050
4	8080	1.0	NO	NONE	...	8 Good	1680

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long \
0	0.0	1955	0.0	98178	47.5112	-122.257
1	400.0	1951	1991.0	98125	47.7210	-122.319
2	0.0	1933	NaN	98028	47.7379	-122.233
3	910.0	1965	0.0	98136	47.5208	-122.393
4	0.0	1987	0.0	98074	47.6168	-122.045

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 21 columns]

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              21597 non-null  int64
1   date            21597 non-null  object
2   price           21597 non-null  float64
```

```

3 bedrooms      21597 non-null int64
4 bathrooms     21597 non-null float64
5 sqft_living    21597 non-null int64
6 sqft_lot       21597 non-null int64
7 floors         21597 non-null float64
8 waterfront     19221 non-null object
9 view           21534 non-null object
10 condition     21597 non-null object
11 grade         21597 non-null object
12 sqft_above    21597 non-null int64
13 sqft_basement 21597 non-null object
14 yr_built      21597 non-null int64
15 yr_renovated  17755 non-null float64
16 zipcode       21597 non-null int64
17 lat           21597 non-null float64
18 long          21597 non-null float64
19 sqft_living15 21597 non-null int64
20 sqft_lot15    21597 non-null int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```

From the above summary: * The data set has 21597 entries * object type columns have to be explored further * date

* waterfront * view

* condition * grade * sqft_basement

0.3 Data preparation

```
[4]: df['date']
```

```

[4]: 0      10/13/2014
     1      12/9/2014
     2      2/25/2015
     3      12/9/2014
     4      2/18/2015
     ...
    21592    5/21/2014
    21593    2/23/2015
    21594    6/23/2014
    21595    1/16/2015
    21596   10/15/2014
Name: date, Length: 21597, dtype: object

```

```

[5]: # Change date column from object to the appropriate data type (datetime)
     # We may later need to re-engineer the date column to derive seasons
     df['date'] = pd.to_datetime(df['date'])

```

```
[6]: print(df['waterfront'].unique())
df['waterfront'].describe()
```

```
[nan 'NO' 'YES']
```

```
[6]: count      19221
unique         2
top           NO
freq         19075
Name: waterfront, dtype: object
```

Waterfront column seems to be a binary column. Having 1 and 0 to represent 'YES' and 'NO' respectively will make analysis much simpler. It's also safe to replace the missing values with the mode; The house not having a view of the waterfront

```
[7]: def waterfront(df):
      df['waterfront'] = df['waterfront'].map(lambda x: 1 if x == 'YES' else 0)
      return df

df = waterfront(df)
df['waterfront'].unique()
```

```
[7]: array([0, 1])
```

```
[8]: print(df['view'].unique())
```

```
['NONE' nan 'GOOD' 'EXCELLENT' 'AVERAGE' 'FAIR']
```

```
[9]: # The null values could be filled with 'NONE'
print(df['view'].isna().value_counts())
```

```
False      21534
True         63
Name: view, dtype: int64
```

```
[10]: print(df['view'].value_counts())
```

```
NONE          19422
AVERAGE       957
GOOD           508
FAIR           330
EXCELLENT      317
Name: view, dtype: int64
```

Using a quantitative variable to represent quality of view from the house would be more effective in EDA

```
[11]: #Function to handle null values and change the view column into int based column
def view(df):
```

```

df['view'].fillna(1,inplace = True)
old = ['NONE', 'AVERAGE', 'FAIR', 'GOOD', 'EXCELLENT']
new = [1,2,3,4,5]

df['view'].replace(old,new,inplace=True)
return df

df = view(df)
df['view'].unique()

```

```
[11]: array([1, 4, 5, 2, 3])
```

```
[12]: df['grade'].unique()
```

```
[12]: array(['7 Average', '6 Low Average', '8 Good', '11 Excellent', '9 Better',
            '5 Fair', '10 Very Good', '12 Luxury', '4 Low', '3 Poor',
            '13 Mansion'], dtype=object)
```

```

[13]: def grade(df):
        old = ['3 Poor', '4 Low', '5 Fair', '6 Low Average', '7 Average', '8 Good', '9_
        ↳Better', '10 Very Good', '11 Excellent', '12 Luxury', '13 Mansion']
        new = [3,4,5,6,7,8,9,10,11,12,13]

        df['grade'].replace(old,new,inplace=True)
        return df
df = grade(df)
df['grade'].unique()

```

```
[13]: array([ 7,  6,  8, 11,  9,  5, 10, 12,  4,  3, 13])
```

```
[14]: df['bathrooms'].unique()
```

```
[14]: array([1.   , 2.25, 3.   , 2.   , 4.5 , 1.5 , 2.5 , 1.75, 2.75, 3.25, 4.   ,
            3.5 , 0.75, 4.75, 5.   , 4.25, 3.75, 1.25, 5.25, 6.   , 0.5 , 5.5 ,
            6.75, 5.75, 8.   , 7.5 , 7.75, 6.25, 6.5 ])
```

- The floating point values in the bathroom column could be indicating the amenities in the bathroom. A whole number could be representing a bathroom with all features, i.e shower, toilet and a sink
- Working with integer values would be more suitable.

```

[15]: # A function to have bathroom columns as int
def bathrooms(df):
    df['bathrooms'] = df['bathrooms'].map(lambda x: int(round(x,0)))
    return df
df = bathrooms(df)
df['bathrooms'].unique()

```

```
[15]: array([1, 2, 3, 4, 5, 6, 0, 7, 8])
```

```
[16]: df['sqft_basement'].unique()
```

```
[16]: array(['0.0', '400.0', '910.0', '1530.0', '?', '730.0', '1700.0', '300.0',  
          '970.0', '760.0', '720.0', '700.0', '820.0', '780.0', '790.0',  
          '330.0', '1620.0', '360.0', '588.0', '1510.0', '410.0', '990.0',  
          '600.0', '560.0', '550.0', '1000.0', '1600.0', '500.0', '1040.0',  
          '880.0', '1010.0', '240.0', '265.0', '290.0', '800.0', '540.0',  
          '710.0', '840.0', '380.0', '770.0', '480.0', '570.0', '1490.0',  
          '620.0', '1250.0', '1270.0', '120.0', '650.0', '180.0', '1130.0',  
          '450.0', '1640.0', '1460.0', '1020.0', '1030.0', '750.0', '640.0',  
          '1070.0', '490.0', '1310.0', '630.0', '2000.0', '390.0', '430.0',  
          '850.0', '210.0', '1430.0', '1950.0', '440.0', '220.0', '1160.0',  
          '860.0', '580.0', '2060.0', '1820.0', '1180.0', '200.0', '1150.0',  
          '1200.0', '680.0', '530.0', '1450.0', '1170.0', '1080.0', '960.0',  
          '280.0', '870.0', '1100.0', '460.0', '1400.0', '660.0', '1220.0',  
          '900.0', '420.0', '1580.0', '1380.0', '475.0', '690.0', '270.0',  
          '350.0', '935.0', '1370.0', '980.0', '1470.0', '160.0', '950.0',  
          '50.0', '740.0', '1780.0', '1900.0', '340.0', '470.0', '370.0',  
          '140.0', '1760.0', '130.0', '520.0', '890.0', '1110.0', '150.0',  
          '1720.0', '810.0', '190.0', '1290.0', '670.0', '1800.0', '1120.0',  
          '1810.0', '60.0', '1050.0', '940.0', '310.0', '930.0', '1390.0',  
          '610.0', '1830.0', '1300.0', '510.0', '1330.0', '1590.0', '920.0',  
          '1320.0', '1420.0', '1240.0', '1960.0', '1560.0', '2020.0',  
          '1190.0', '2110.0', '1280.0', '250.0', '2390.0', '1230.0', '170.0',  
          '830.0', '1260.0', '1410.0', '1340.0', '590.0', '1500.0', '1140.0',  
          '260.0', '100.0', '320.0', '1480.0', '1060.0', '1284.0', '1670.0',  
          '1350.0', '2570.0', '1090.0', '110.0', '2500.0', '90.0', '1940.0',  
          '1550.0', '2350.0', '2490.0', '1481.0', '1360.0', '1135.0',  
          '1520.0', '1850.0', '1660.0', '2130.0', '2600.0', '1690.0',  
          '243.0', '1210.0', '1024.0', '1798.0', '1610.0', '1440.0',  
          '1570.0', '1650.0', '704.0', '1910.0', '1630.0', '2360.0',  
          '1852.0', '2090.0', '2400.0', '1790.0', '2150.0', '230.0', '70.0',  
          '1680.0', '2100.0', '3000.0', '1870.0', '1710.0', '2030.0',  
          '875.0', '1540.0', '2850.0', '2170.0', '506.0', '906.0', '145.0',  
          '2040.0', '784.0', '1750.0', '374.0', '518.0', '2720.0', '2730.0',  
          '1840.0', '3480.0', '2160.0', '1920.0', '2330.0', '1860.0',  
          '2050.0', '4820.0', '1913.0', '80.0', '2010.0', '3260.0', '2200.0',  
          '415.0', '1730.0', '652.0', '2196.0', '1930.0', '515.0', '40.0',  
          '2080.0', '2580.0', '1548.0', '1740.0', '235.0', '861.0', '1890.0',  
          '2220.0', '792.0', '2070.0', '4130.0', '2250.0', '2240.0',  
          '1990.0', '768.0', '2550.0', '435.0', '1008.0', '2300.0', '2610.0',  
          '666.0', '3500.0', '172.0', '1816.0', '2190.0', '1245.0', '1525.0',  
          '1880.0', '862.0', '946.0', '1281.0', '414.0', '2180.0', '276.0',  
          '1248.0', '602.0', '516.0', '176.0', '225.0', '1275.0', '266.0',  
          '283.0', '65.0', '2310.0', '10.0', '1770.0', '2120.0', '295.0',
```

```
'207.0', '915.0', '556.0', '417.0', '143.0', '508.0', '2810.0',
'20.0', '274.0', '248.0'], dtype=object)
```

```
[17]: df['sqft_basement'].value_counts()
```

```
[17]: 0.0      12826
      ?        454
      600.0    217
      500.0    209
      700.0    208
      ...
      1920.0     1
      3480.0     1
      2730.0     1
      2720.0     1
      248.0      1
      Name: sqft_basement, Length: 304, dtype: int64
```

```
[18]: df['sqft_basement'].describe()
```

```
[18]: count      21597
      unique      304
      top         0.0
      freq      12826
      Name: sqft_basement, dtype: object
```

- From the summary statistics it seems many houses don't have a basement hence 0 sqft. The '?' is a place holder for 454 houses in which the basement status is not indicated
- Having the basement status as a binary option could provide better insights as compared to the size.

```
[19]: def basement(df):
      df['sqft_basement'] = df['sqft_basement'].map(lambda x : float(x.replace('?',
      ↪', '0'))))
      df['basement'] = df['sqft_basement'].map(lambda x: 1 if x > 0 else 0)
      df['sqft_basement'] = df['basement']

      return df

df = basement(df)
df['basement'].unique()
```

```
[19]: array([0, 1])
```

```
[20]: # Let's check for duplicates and remove them if any
      df[df.duplicated()==True].shape
```

```
[20]: (0, 22)
```

```
[21]: # Determinig missing values
df.isnull().sum()
```


```
[21]: id                0
      date             0
      price            0
      bedrooms         0
      bathrooms        0
      sqft_living       0
      sqft_lot          0
      floors            0
      waterfront        0
      view              0
      condition         0
      grade             0
      sqft_above        0
      sqft_basement     0
      yr_built          0
      yr_renovated      3842
      zipcode           0
      lat               0
      long              0
      sqft_living15     0
      sqft_lot15        0
      basement          0
      dtype: int64
```

- There seems to be no duplicated rows across the 22 columns
- Most of the null values were handled in the initial feature engineering apart from yr_renovated

```
[22]: df['yr_renovated'].dropna(inplace=True)
```

0.3.1 Data Analysis

```
[23]: sns.barplot(x='grade',y='price',data=df)
      plt.title("Grade vs Price")
      plt.show()
```

- The grade of the house refers to the general quality of the house with 3 being poor quality and 13 being more than luxurious.
- Renovations are likely to improve the grade and general condition of the house hence contributing to the overall price.

```
[24]: df['bedrooms'].value_counts()
```

```
[24]: 3      9824
      4      6882
      2      2760
      5      1601
      6       272
      1       196
      7        38
      8        13
      9         6
     10         3
     11         1
     33         1
      Name: bedrooms, dtype: int64
```

```
[25]: df[df['bedrooms'] > 10 ]
```

```
[25]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
8748	1773100755	2014-08-21	520000.0	11	3	3000	
15856	2402100895	2014-06-25	640000.0	33	2	1620	

	sqft_lot	floors	waterfront	view	...	sqft_above	sqft_basement	\
8748	4960	2.0	0	1	...	2400	1	
15856	6000	1.0	0	1	...	1040	1	

	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	\
8748	1918	1999.0	98106	47.5560	-122.363	1420	
15856	1947	0.0	98103	47.6878	-122.331	1330	


	sqft_lot15	basement
8748	4960	1
15856	4700	1


```
[2 rows x 22 columns]
```

The two houses with 11 and 33 bedrooms have a significantly lower ratio of bathrooms. It's highly likely to be an anomaly and should therefore be dropped.


```
[26]: outliers = df[df['bedrooms'] > 10 ].index
df.drop(outliers, inplace=True)
```

```
[27]: # Bedroom distribution
counts = df["bedrooms"].value_counts()
sns.barplot(x=counts.index,y=counts.values,palette=("Set2"))
plt.xlabel("No. of Bedrooms")
plt.title("Bedroom distribution")
plt.show()
```




A typical household has 3 bedrooms followed by 4 bedroom and 3 bedroom houses.

```
[28]: sns.barplot(x='view',y='price',data=df)  
plt.title('View vs Price');
```



```
[29]: sns.boxplot(x = df['waterfront'], y = df['price'])  
plt.xticks(np.arange(2), ('Not view of waterfront', 'Waterfront view'))  
plt.title("Boxplot of waterfront feature vs. price")  
plt.show()
```




```
[30]: waterfront = df[df['waterfront'] == 1]['price'].mean()
non_waterfront = df[df['waterfront'] == 0]['price'].mean()
print(f"A house with a waterfront has an average price of USD_
↳{round(waterfront,2)}")
print(f"Houses without a waterfront have an average price of USD_
↳{round(non_waterfront,2)}")
```

A house with a waterfront has an average price of USD 1717214.73
Houses without a waterfront have an average price of USD 532281.77

```
[31]: df['sale_year'] = df.apply(lambda x: x.date.year, axis=1)
df['sale_year'].value_counts()
```

```
[31]: 2014    14620
      2015     6975
      Name: sale_year, dtype: int64
```

```
[32]: sns.countplot(x = df['sale_year'])
plt.title('House Sales in 2014-2015')
plt.show()
```

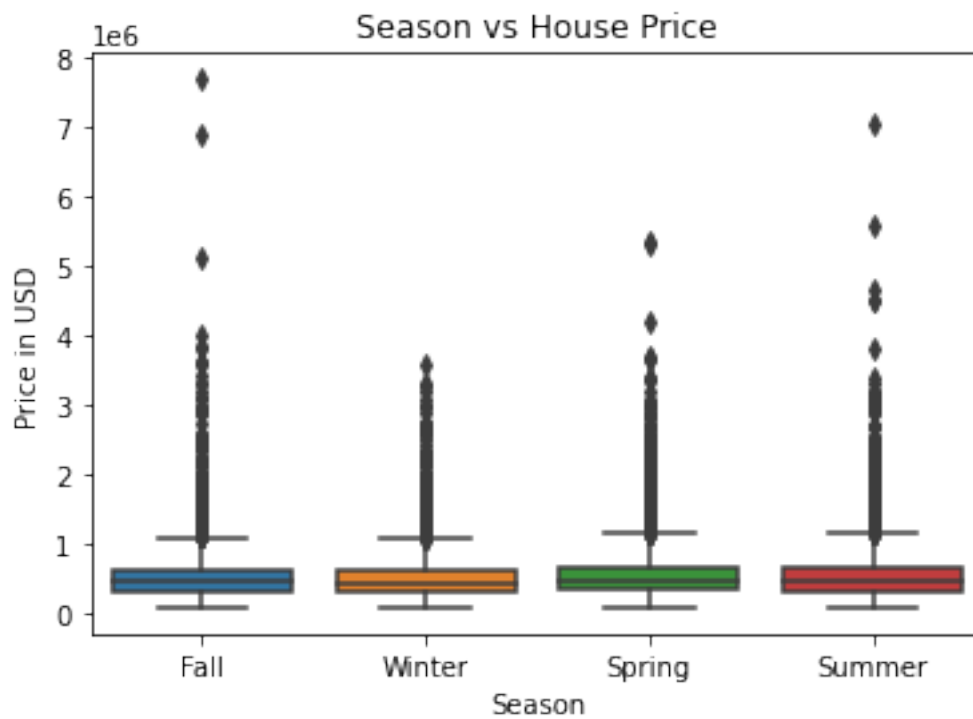


```
[33]: # A function to determine season from the date column
def monthToSeason(df):
    df['sale_month'] = df.apply(lambda x: x.date.month, axis=1)
    seasons={1: 'Winter',2: 'Winter',3: 'Spring',4: 'Spring',5: 'Spring',6: 'Summer',
              7: 'Summer',8: 'Summer',9: 'Fall',10: 'Fall',11: 'Fall',12: 'Winter'}
    df['season']=df.sale_month.map(seasons)
    return df


df = monthToSeason(df)
df['season'].unique()
```

```
[33]: array(['Fall', 'Winter', 'Spring', 'Summer'], dtype=object)
```

```
[34]: sns.boxplot(x = df['season'], y = df['price'], data = df)
plt.xlabel('Season')
plt.ylabel('Price in USD')
plt.title('Season vs House Price')
plt.show()
```



```
[35]: sns.countplot(x=df['season'])  
plt.title('House Sale Across Seasons')  
plt.show()
```



- There's an interesting observation in house sales across the years, almost twice as many houses were sold in 2014 as compared to 2015.
- The season has little to no impact on the overall price of the house.
- Many houses are sold/bought in Spring and Summer.
- Gearing our market campaigns during Fall and Winter could get the ball rolling and complete purchase/sale in Spring and Summer

0.3.2 Modelling

```
[36]: df.sort_values("price", ascending = False).head()
```

```
[36]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
7245	6762700020	2014-10-13	7700000.0	6	8	12050	
3910	9808700762	2014-06-11	7060000.0	5	4	10040	
9245	9208900037	2014-09-19	6890000.0	6	8	9890	
4407	2470100110	2014-08-04	5570000.0	5	6	9200	
1446	8907500070	2015-04-13	5350000.0	5	5	8000	

	sqft_lot	floors	waterfront	view	...	yr_renovated	zipcode	lat	\
7245	27600	2.5	0	4	...	1987.0	98102	47.6298	
3910	37325	2.0	1	2	...	2001.0	98004	47.6500	
9245	31374	2.0	0	5	...	0.0	98039	47.6305	
4407	35069	2.0	0	1	...	NaN	98039	47.6289	
1446	23985	2.0	0	5	...	0.0	98004	47.6232	

	long	sqft_living15	sqft_lot15	basement	sale_year	sale_month	\
7245	-122.323	3940	8800	1	2014	10	
3910	-122.214	3930	25449	1	2014	6	
9245	-122.240	4540	42730	1	2014	9	
4407	-122.233	3560	24345	1	2014	8	
1446	-122.220	4600	21750	1	2015	4	

	season
7245	Fall
3910	Summer
9245	Fall
4407	Summer
1446	Spring

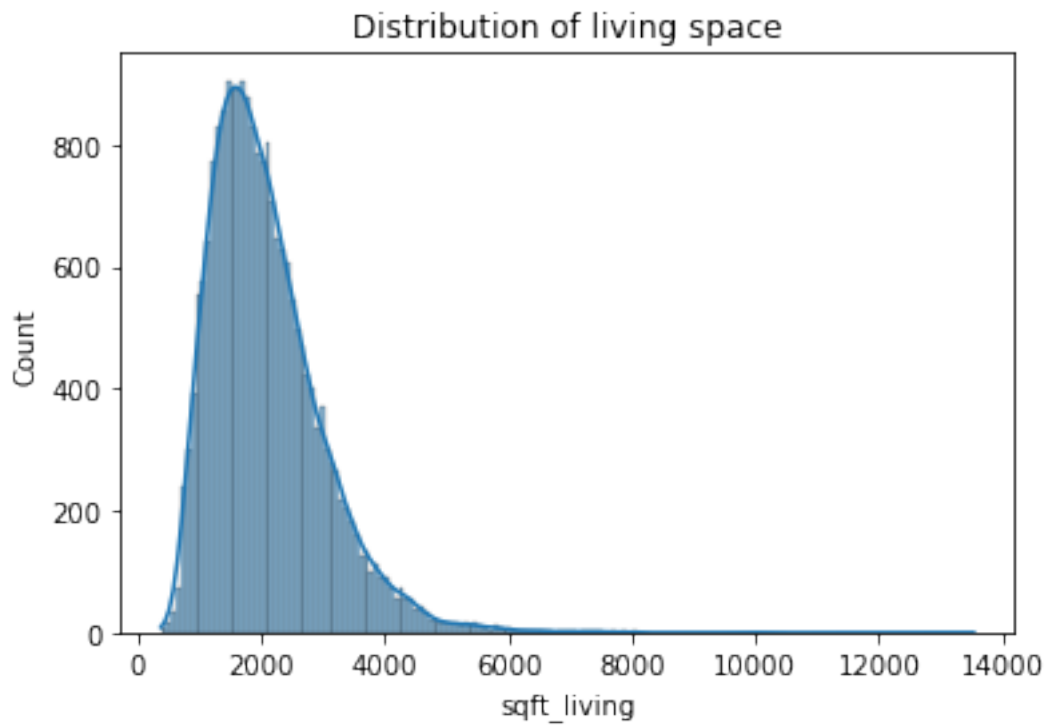
[5 rows x 25 columns]

```
[37]: #Explore the relationship between price and other columns
df.corr('pearson')['price']
```


```
[37]: id                -0.016765
price                 1.000000
bedrooms              0.316504
bathrooms             0.519652
sqft_living           0.701948
sqft_lot              0.089879
floors                0.256828
waterfront            0.264308
view                  0.394713
grade                 0.667967
sqft_above            0.605401
sqft_basement         0.178264
yr_built              0.053964
yr_renovated          0.129702
zipcode              -0.053408
lat                   0.306687
long                  0.022045
sqft_living15         0.585274
sqft_lot15            0.082848
basement              0.178264
sale_year             0.003734
sale_month            -0.009925
Name: price, dtype: float64
```

```
[38]: # living space Distribution
sns.histplot(df['sqft_living'], kde=True);
plt.title("Distribution of living space")
```

```
plt.show()
```



```
[39]: # The 'sqft_living' column has the highest relationship with price
df.plot(x='sqft_living',y='price', kind='scatter')
plt.title("Living space vs Price")
plt.show()
```



- From the two distributions, most of the houses have a living space within 6000 sqft and tend to cost a figure below \$4,000,000

```
[40]: train_data, test_data = train_test_split(df, train_size=0.8, random_state=3)
reg = LinearRegression()
x_train = np.array(train_data['sqft_living']).reshape(-1, 1)
y_train = np.array(train_data['price']).reshape(-1, 1)
reg.fit(x_train, y_train)

x_test = np.array(test_data['sqft_living']).reshape(-1, 1)
y_test = np.array(test_data['price']).reshape(-1, 1)
pred = reg.predict(x_test)
```

```
[41]: plt.subplots(figsize = (9, 7))
plt.scatter(x_test, y_test, color= 'blue', label = 'data')
plt.plot(x_test, reg.predict(x_test), color = 'red', label = ' Predicted_
↳ Regression line')
plt.xlabel('Living Space (sqft)')
plt.ylabel('price')
plt.legend()
plt.gca().spines['right'].set_visible(False)
plt.gca().spines['right'].set_visible(False)

#Model summary
```

```

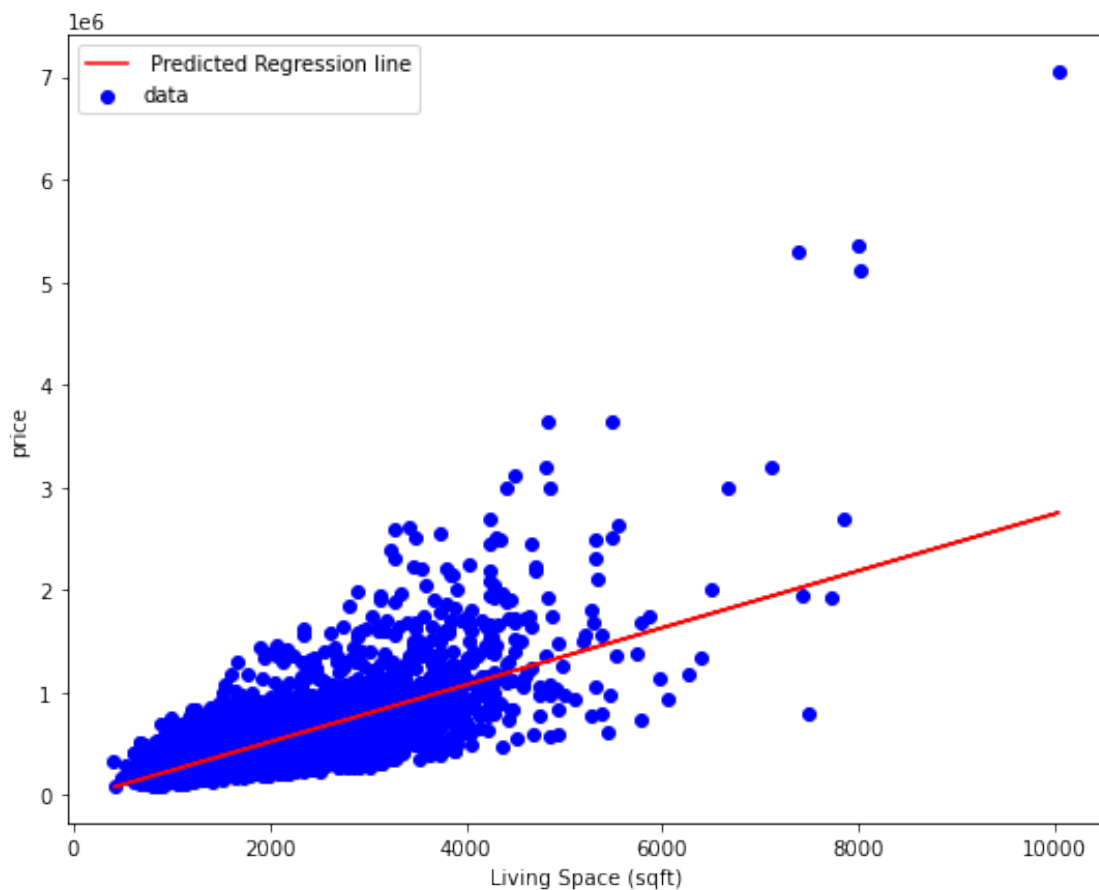
mean_squared_error=metrics.mean_squared_error(y_test,pred)
print('Sqaured mean error', round(np.sqrt(mean_squared_error),2))
print('R squared training',round(reg.score(x_train,y_train),3))
print('R sqaured testing',round(reg.score(x_test,y_test),3) )
print('intercept',reg.intercept_)
print('coefficient',reg.coef_)

```

```

Sqaured mean error 273789.66
R squared training 0.494
R sqaured testing 0.488
intercept [-38017.20149875]
coefficient [[277.90934404]]

```



- After living space, grade of the house is the next variable to have a strong correlation with price.

```

[42]: train_data,test_data = train_test_split(df,train_size=0.8,random_state=3)
reg = LinearRegression()
x_train = np.array(train_data['grade']).reshape(-1,1)

```

```

y_train = np.array(train_data['price']).reshape(-1,1)
reg.fit(x_train,y_train)

x_test = np.array(test_data['grade']).reshape(-1,1)
y_test = np.array(test_data['price']).reshape(-1,1)
pred = reg.predict(x_test)

print("Model 1")
mean_squared_error=metrics.mean_squared_error(y_test,pred)
print('squared mean error',round(np.sqrt(mean_squared_error),2))
print('R squared training',round(reg.score(x_train,y_train),3))
print('R squared testing',round(reg.score(x_test,y_test),3))
print('intercept',reg.intercept_)
print('coefficient',reg.coef_)

```

```

Model 1
squared mean error 288645.26
R squared training 0.45
R squared testing 0.431
intercept [-1050417.80366832]
coefficient [[207678.58717913]]


```

0.3.3 Multiple regression

```

[43]: fig,ax=plt.subplots(3,1,figsize=(15,10))
sns.boxplot(x=train_data['grade'],y=train_data['price'],ax=ax[0])
sns.boxplot(x=train_data['bedrooms'],y=train_data['price'],ax=ax[1])
sns.boxplot(x=train_data['bathrooms'],y=train_data['price'],ax=ax[2])
plt.show()

```



```
[44]: model_2 = ['bedrooms', 'grade', 'sqft_living', 'sqft_above']
reg=LinearRegression()
reg.fit(train_data[model_2], train_data['price'])
pred=reg.predict(test_data[model_2])
print('Model 2')
mean_squared_error=metrics.mean_squared_error(y_test, pred)
print('mean squared error(MSE)', round(np.sqrt(mean_squared_error), 2))
print('R squared training', round(reg.
    ↳score(train_data[model_2], train_data['price']), 3))
print('R squared testing', round(reg.
    ↳score(test_data[model_2], test_data['price']), 3))
print('Intercept: ', reg.intercept_)
print('Coefficient:', reg.coef_)
```

Model 2

mean squared error(MSE) 260437.56

R squared training 0.554

R squared testing 0.537

Intercept: -513600.3368260097

Coefficient: [-4.73070108e+04 1.04713234e+05 2.68444581e+02 -8.22739427e+01]

```
[45]: model_3 =
    ↪ ['basement', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'view', 'g
reg = LinearRegression()
reg.fit(train_data[model_3], train_data['price'])
pred = reg.predict(test_data[model_3])

print('Model 3')
mean_squared_error = metrics.mean_squared_error(y_test, pred)
print('Mean Squared Error (MSE) ', round(np.sqrt(mean_squared_error), 2))
print('R-squared (training) ', round(reg.score(train_data[model_3],
    ↪ train_data['price']), 3))
print('R-squared (testing) ', round(reg.score(test_data[model_3],
    ↪ test_data['price']), 3))
print('Intercept: ', reg.intercept_)
print('Coefficient:', reg.coef_)
```

```
Model 3
Mean Squared Error (MSE)  222830.8
R-squared (training)  0.659
R-squared (testing)  0.661
Intercept:  -31881653.514646027
Coefficient: [ 2.13948976e+03 -3.24404951e+04  2.20573040e+04  1.77897381e+02
 -1.94735194e-01 -3.82042849e+04  5.25017963e+05  7.83235763e+04
 8.10187710e+04  1.27264594e+01  6.60923385e+05]
```

0.3.4 Model summary

- Our third model better as compared to our baseline model.
- It was a better predictor of price as it had a lower error value and a higher R-Squared value

0.3.5 Conclusions

- An increase in the number of **bedrooms** and **bathrooms** increases the price of the house. Houses with a good view also tend to fetch high prices
- Grade is a good indicator of price. Renovations will therefore contribute to price as it improves the house condition and grade to some extent.
- Houses with a waterfront cost thrice as much as houses without. This are high end properties that should have customised marketing campaign.
- The season or time of the year was a prospect in determining price but it turns out there's no significant contribution. However, the distribution gives some insight on timing of our marketing campaigns.